

Predicting Wage by Club & Player Attributes

Justo Garcia Miranda Heyer James Kistner Cady Stringer

Dataset

- ❑ FIFA '19 Player Statistics
 - ❑ 18,000 players
 - ❑ 90 attributes



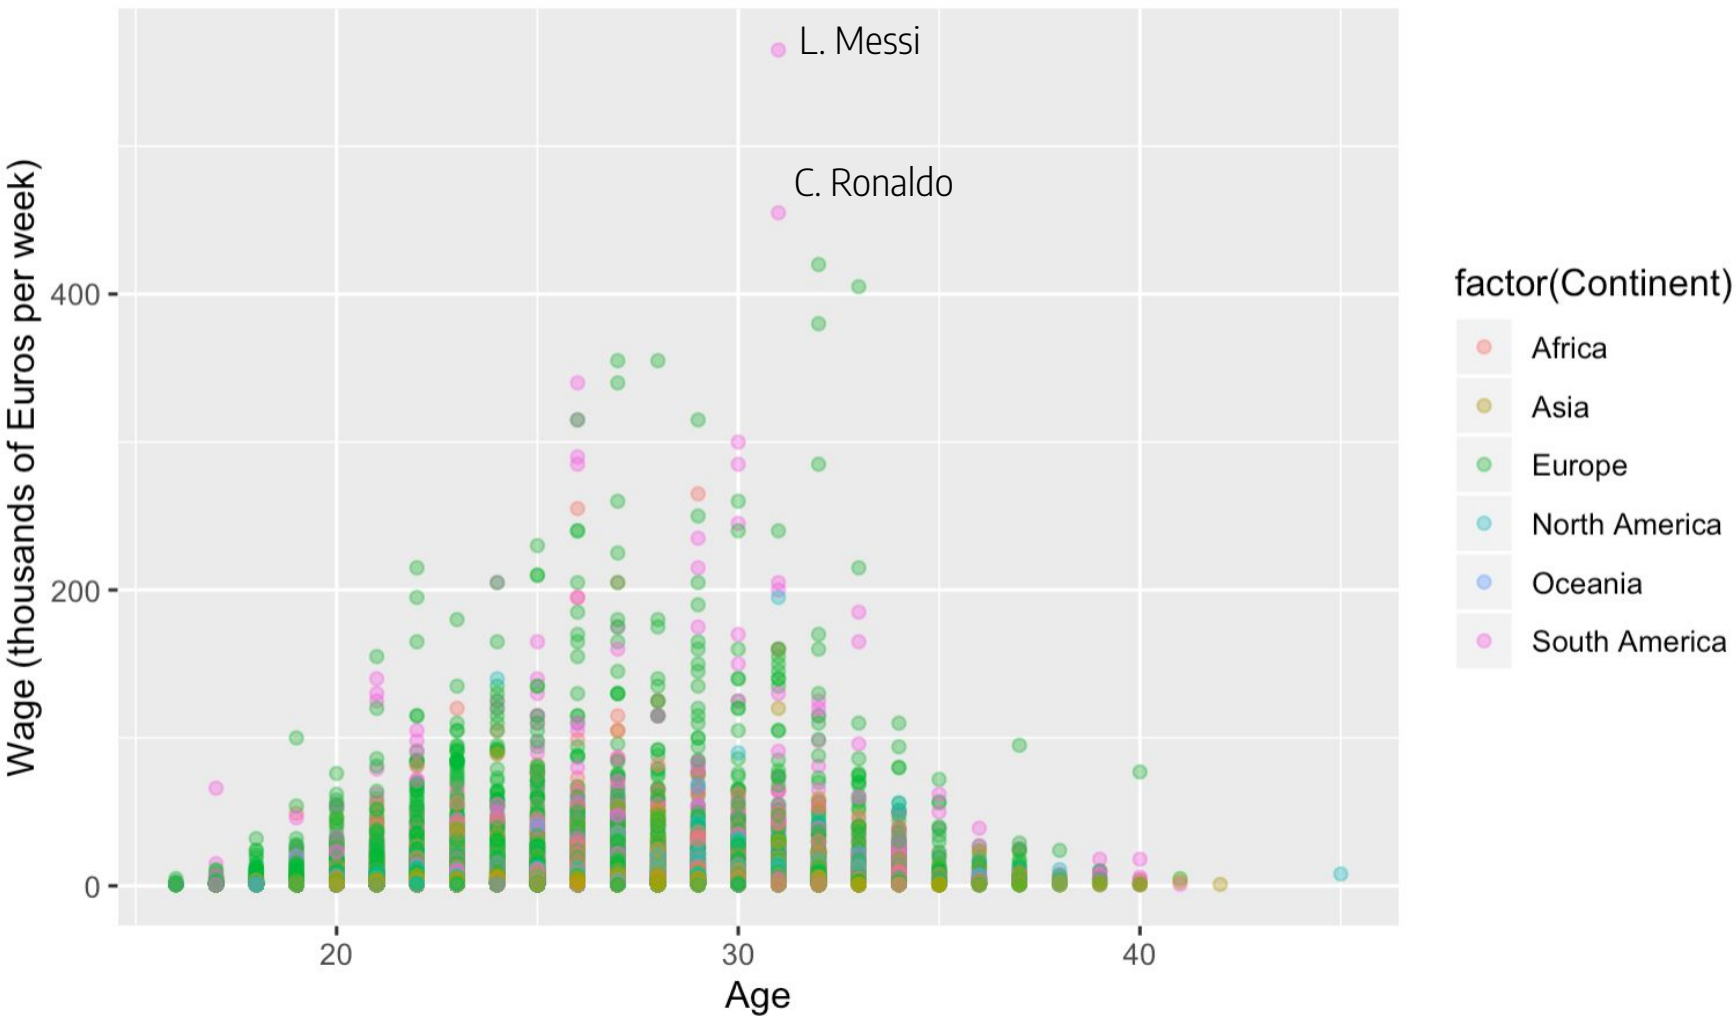
Business Use Case

- ▶ Soccer team management can use our model to predict what a player is worth *before* negotiating a contract
 - ▶ Better allocate market cap (all the money they have for player salaries) to players
 - ▶ Helps them decide when to splurge and when to save




Data Exploration

Player Age vs Wage colored by Continent




Player Wage vs Overall Skill Rating



Club

● Arsenal	● Liverpool
● AS Monaco	● Manchester City
● Atlético Madrid	● Manchester United
● Borussia Dortmund	● Newcastle United
● Burnley	● Rayo Vallecano
● Cardiff City	● RC Celta
● CD Leganés	● Real Madrid
● Chelsea	● Southampton
● Eintracht Frankfurt	● Tottenham Hotspur
● Empoli	● TSG 1899 Hoffenheim
● Everton	● Valencia CF
● FC Barcelona	● Wolverhampton Wanderers
● Fortuna Düsseldorf	● Other
● Frosinone	

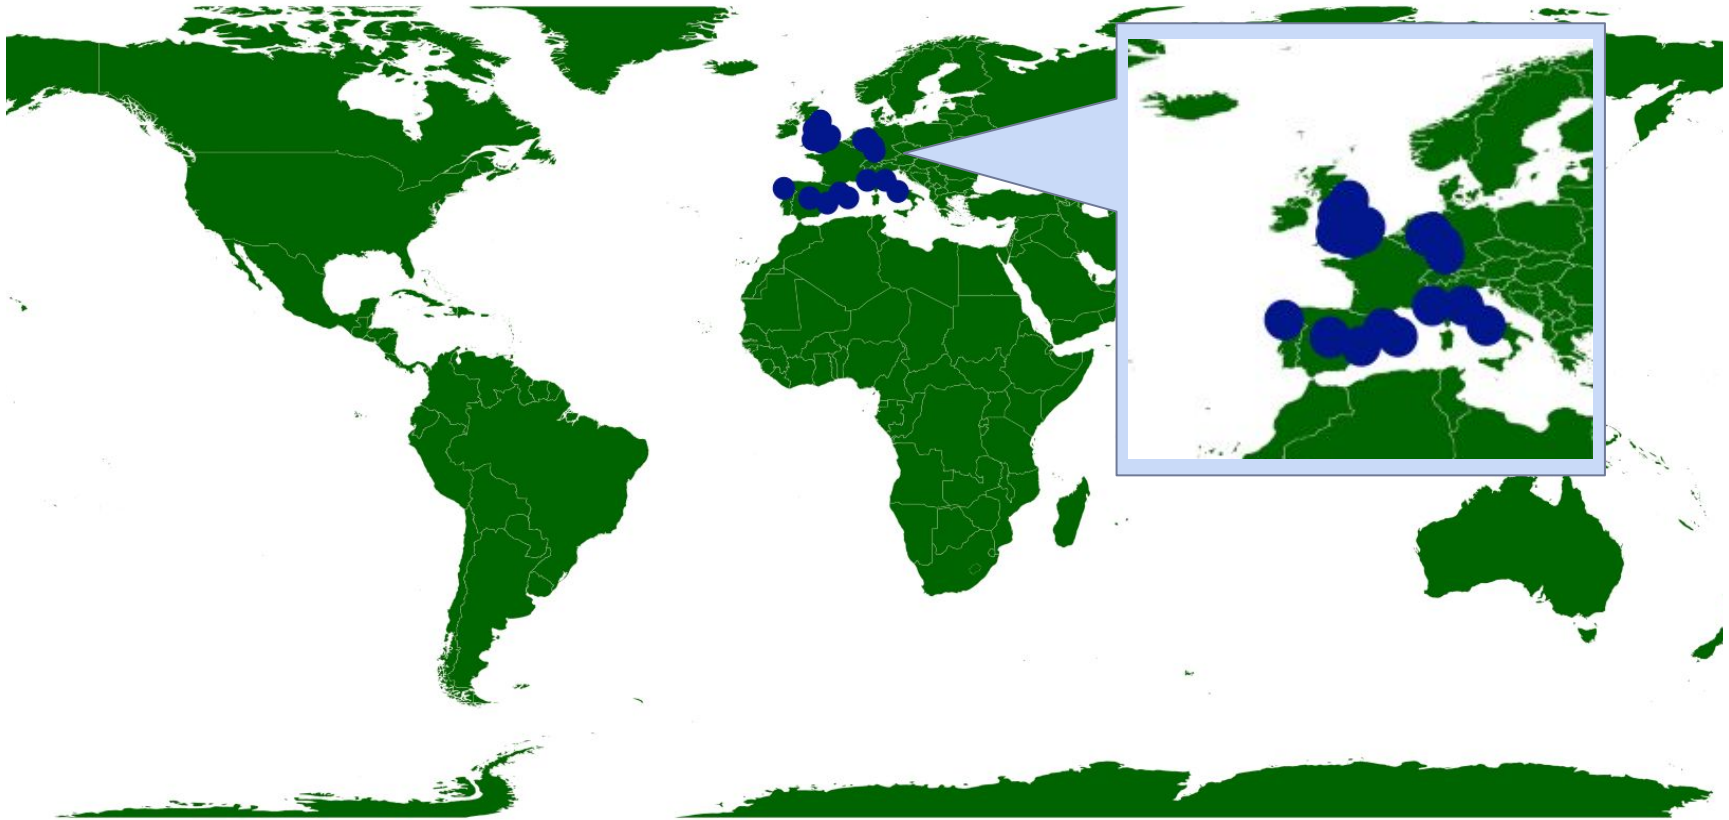
Wage Density by Club



Club



Top 28 Club Locations





Data Cleaning

Data Cleaning

- ▷ Remove unneeded columns
 - ▶ Row #, ID #, Value, Release Clause, Name, Photo, Flag, Club Logo, Special, Body Type, Real Face, Jersey #, Loaned From, Joined Date
- ▷ Removed symbols/letters in numeric columns
 - ▶ Wage & weight
 - ▶ Height to inches function removes apostrophe, replaces it with decimal and changes it to inches 5'11 to 71 inches
- ▷ Factoring
 - ▶ Over 200 countries, so we changed countries to continents by a right_join
 - ▶ Factored clubs into top 28 & other



Models

Linear Regression, Linear Regression with Lasso
Variable Selection, Elastic Net, Random Forest, Neural
Network


Multiple Linear Regression

```
fifa_lm_mod <- lm(Wage ~ .,  
                  fifa_train)  
preds_DF_lm <- data.frame(fifa_lm_preds = predict(fifa_lm_mod, newdata=fifa_test))  
postResample(preds_DF_lm$fifa_lm_preds, fifa_test$Wage)
```


Linear Model Diagnostics


- ▷ Better than expected!
- ▷ R^2 : 0.66
- ▷ MAE: 6.4
- ▷ RMSE: 13.8

Heteroskedasticity

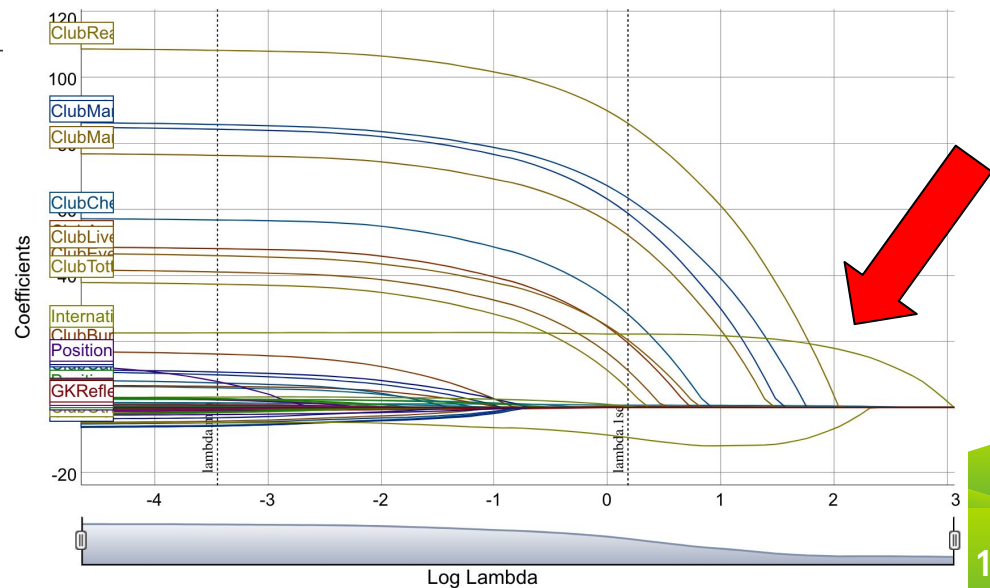


Elastic Net: Picking our Optimal Alpha





- *Clubs are some of the last variables to be set to zero*
- *International reputation last variable set to zero*



Elastic Net Model Diagnostics

- ▷ R^2 : 0.63
- ▷ MAE: 6.3
- ▷ RMSE: 14.5

Lasso Model


```
lasso_mod_fifa <- cv.glmnet(Wage ~ .,  
                           data = fifa_train,  
                           alpha = 1,  
                           nfolds = 10)  
  
fifa_lm_mod_lassopicks <- lm(Wage ~ International.Reputation + Overall + Potential + Club,  
                             fifa_train)  
  
preds_DF_lm_lassopicks <- data.frame(fifa_lm_preds = predict(fifa_lm_mod_lassopicks, newdata=fifa_test))  
postResample(preds_DF_lm_lassopicks$fifa_lm_preds, fifa_test$Wage)
```

Lasso Model

- ▷ Lambda min picked 96 variables
- ▷ Lambda 1se picked 8 variables:
 - ▶ International.Reputation
 - ▶ Overall
 - ▶ Potential
 - ▶ 4 clubs and “Other” club factor

Lasso Model Diagnostics

- ▷ R^2 : 0.65
- ▷ MAE: 6.3
- ▷ RMSE: 14



Post Lasso Estimator

- ▷ We used the variable selections from Lasso (λ_{1se}) to make a linear model to see how it performs compared to one with all variables

Post Lasso Estimator


- ▷ LM with Lasso variables
 - ▷ R²: 0.649
 - ▷ RMSE: 14.04
 - ▷ MAE: 6.27
 - ▷ Those 8 variables do almost as well as a model using all the variables
- ▷ Regular LM
 - ▷ R²: 0.659
 - ▷ RMSE: 12.84
 - ▷ MAE: 6.39

Random Forest

- ▶ Used everything but Club because it had too many categories
- ▶ Chose 100 trees to minimize error

```
fifa_train_noclub <- fifa_train %>% select(-"Club")  
  
rf_fifa_fit <- randomForest(Wage ~.,  
                           data = fifa_train_noclub,  
                           type = regression,  
                           ntree = 100,  
                           importance = TRUE,  
                           localImp = TRUE  
                           )
```


rf_fifa_fit




rf_fifa_fit


Top 5
variables




Overall
Potential
International.Reputation
Reactions
Position
Age
StandingTackle
Marking
Continent
Volleys
Finishing
Stamina
SlidingTackle
Interceptions
BallControl
Strength
LongPassing
Aggression
Height
SprintSpeed
ShortPassing
GKReflexes
Composure
FKAccuracy
Acceleration
Weight
GKDiving
Vision
Dribbling
Agility




Overall
International.Reputation
Reactions
Potential
BallControl
Position
ShortPassing
Composure
Age
Dribbling
StandingTackle
LongShots
SlidingTackle
Vision
Positioning
Finishing
Interceptions
ShotPower
Work.Rate
GKPositioning
Jumping
Stamina
Aggression
FKAccuracy
GKHandling
Marking
Volleys
Curve
LongPassing
HeadingAccuracy



Prediction of the forest for different values of Overall and Potential



Prediction of the forest for different values of Overall and International.Reputation

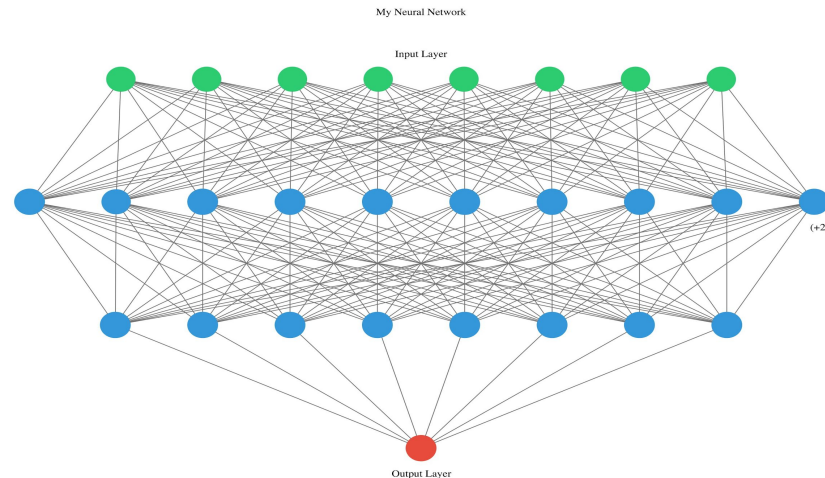


Random Forest Model Diagnostics

- ▷ R^2 : 0.79
- ▷ MAE: 4.5
- ▷ RMSE: 11.3

Neural Network

- ▷ Feed input into network
- ▷ Make prediction
- ▷ Update weights through backpropagation



Neural Network

```
import pandas as pd
import keras
import numpy as np
import matplotlib.pyplot as plt
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout
x_train = pd.read_csv("datasets/fifa_xtrain.csv")
x_test = pd.read_csv("datasets/fifa_xtest.csv")
y_train = pd.read_csv("datasets/fifa_ytrain.csv")
y_test = pd.read_csv("datasets/fifa_ytest.csv")
```

Using TensorFlow backend.

```
y_train = (y_train - y_train.mean()) / y_train.std()
x_train = (x_train - x_train.mean()) / x_train.std()
y_test = (y_test - y_test.mean()) / y_test.std()
x_test = (x_test - x_test.mean()) / x_test.std()
```


```
print(x_train.shape)
print(x_test.shape)
```

```
(9737, 42)
(4173, 42)
```

```
model = Sequential()
model.add(Dense(42, input_dim = 42, kernel_initializer = 'normal', activation='relu'))
model.add(Dense(21, kernel_initializer='normal', activation = 'relu'))
model.add(Dense(10, kernel_initializer='normal', activation = 'relu'))
model.add(Dropout(0.5))
model.add(Dense(5, kernel_initializer='normal', activation = 'relu'))
model.add(Dense(1, kernel_initializer='normal'))
model.compile(loss='mean_squared_error', optimizer='adam', metrics = ['mape'])
```

```
history = model.fit(
    x_train, y_train, # training data to learn from
    batch_size= 107, # size of batches
    epochs= 75, # how many iterations we train for
    verbose=1, # type of logging
    validation_data=(x_test, y_test))
```

Neural Network Diagnostics



▷ R2: 0.74

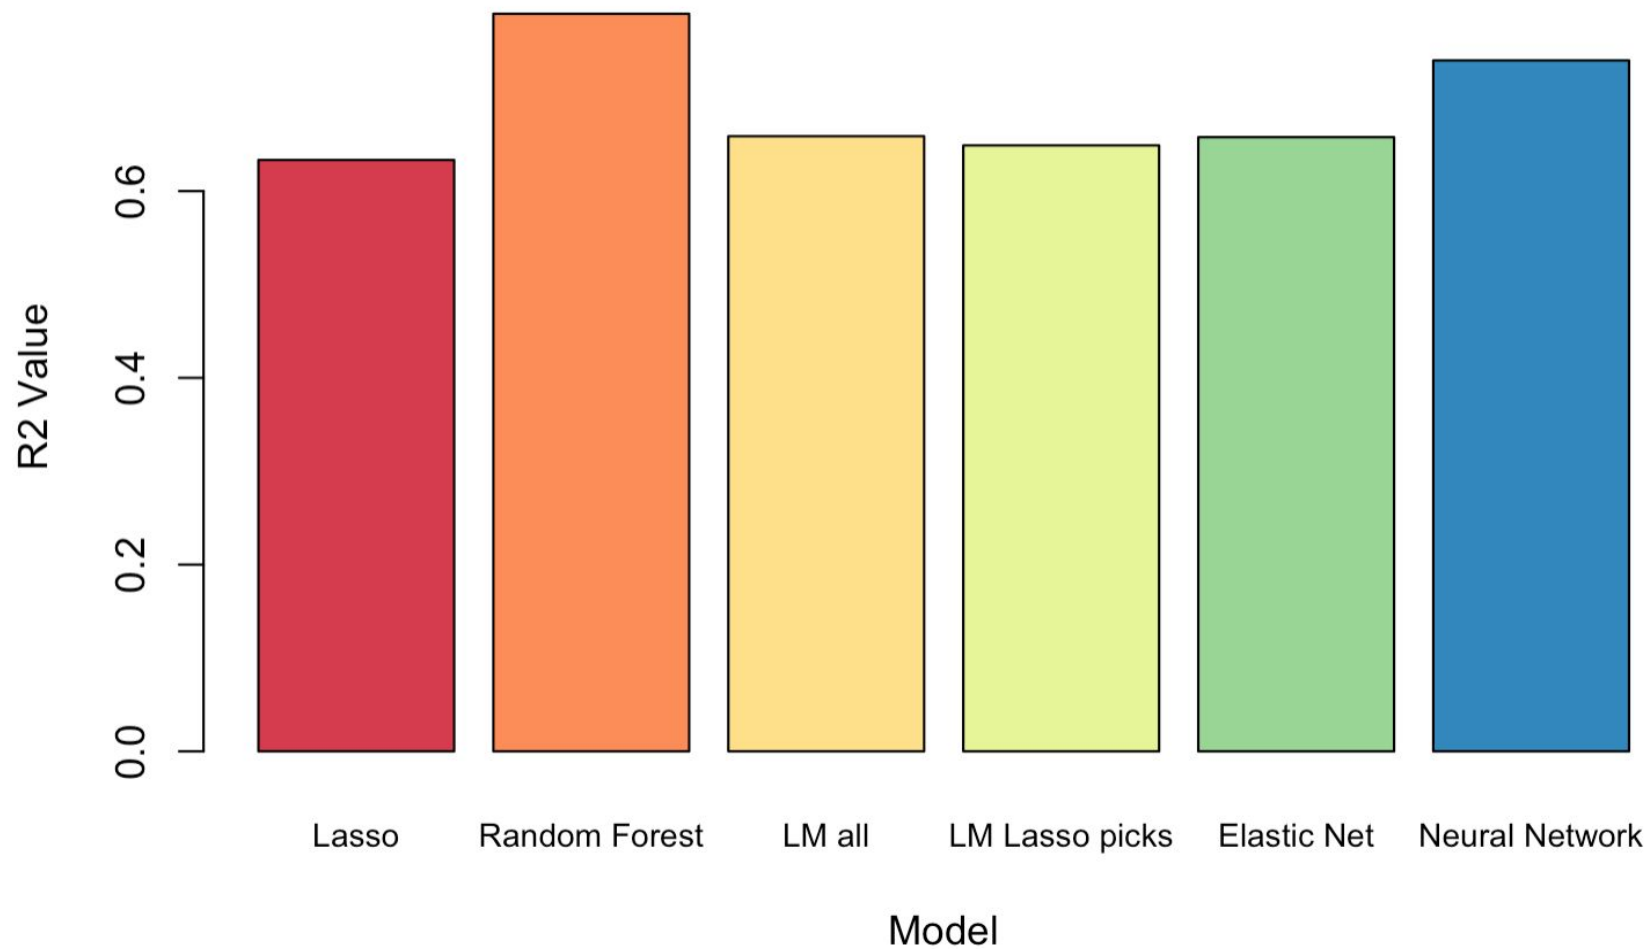


Results

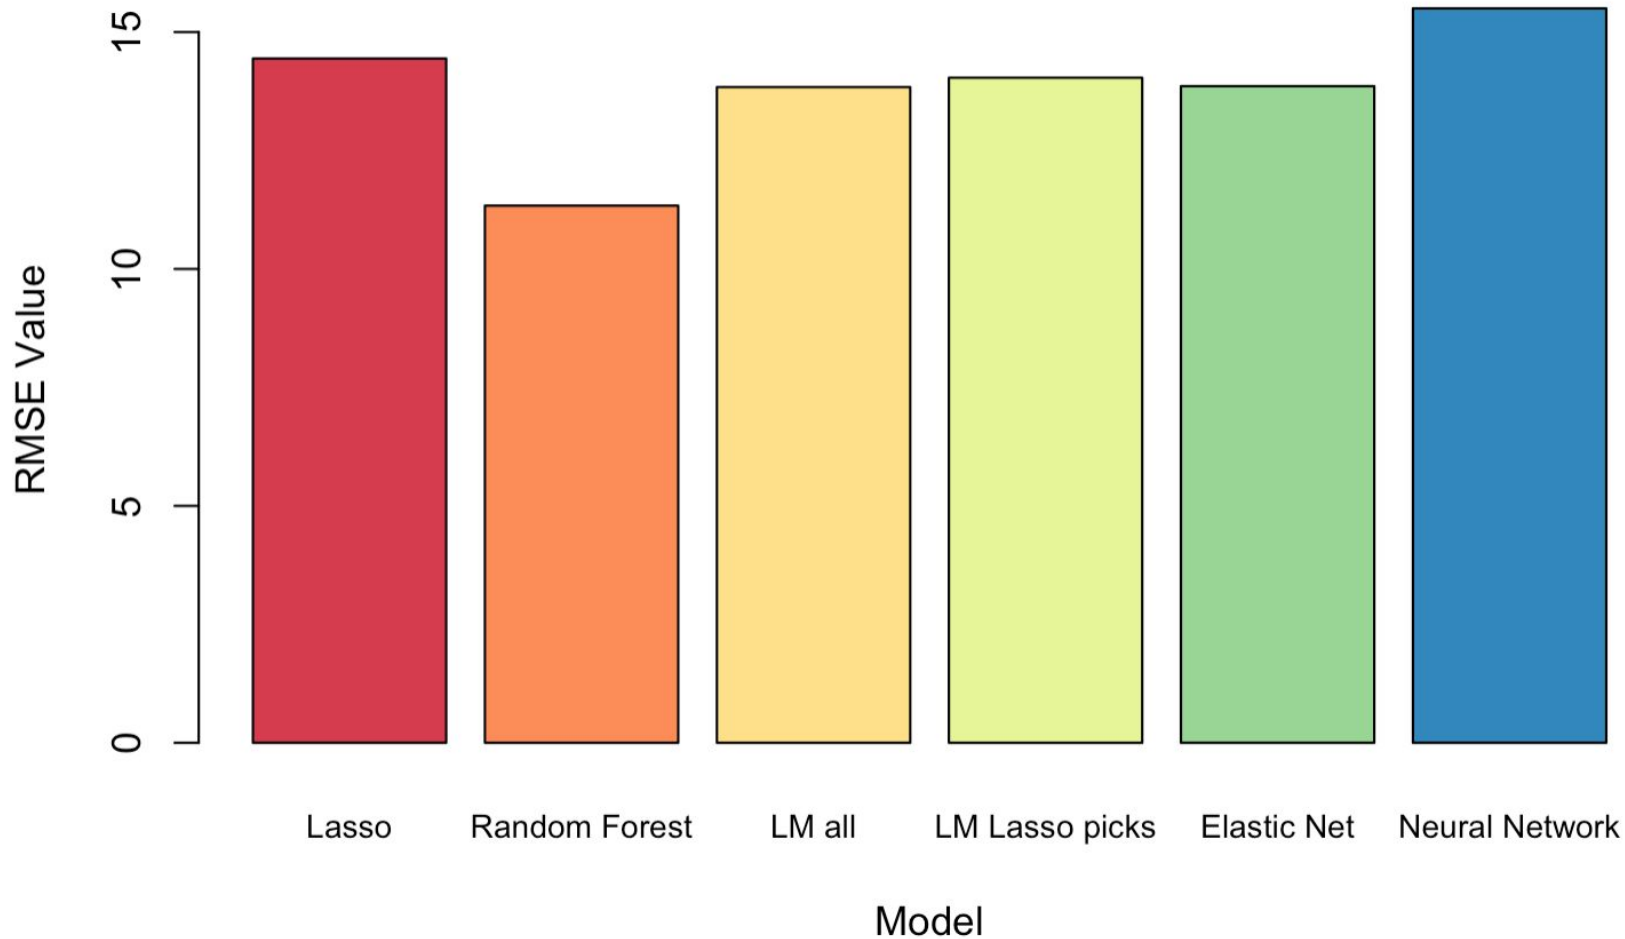
Results

- ▷ Most important variables across all models:
 - ▶ Overall skill rating
 - ▶ Potential rating
 - ▶ Club
 - ▶ International Reputation
 - ▶ Reactions
 - ▶ Position

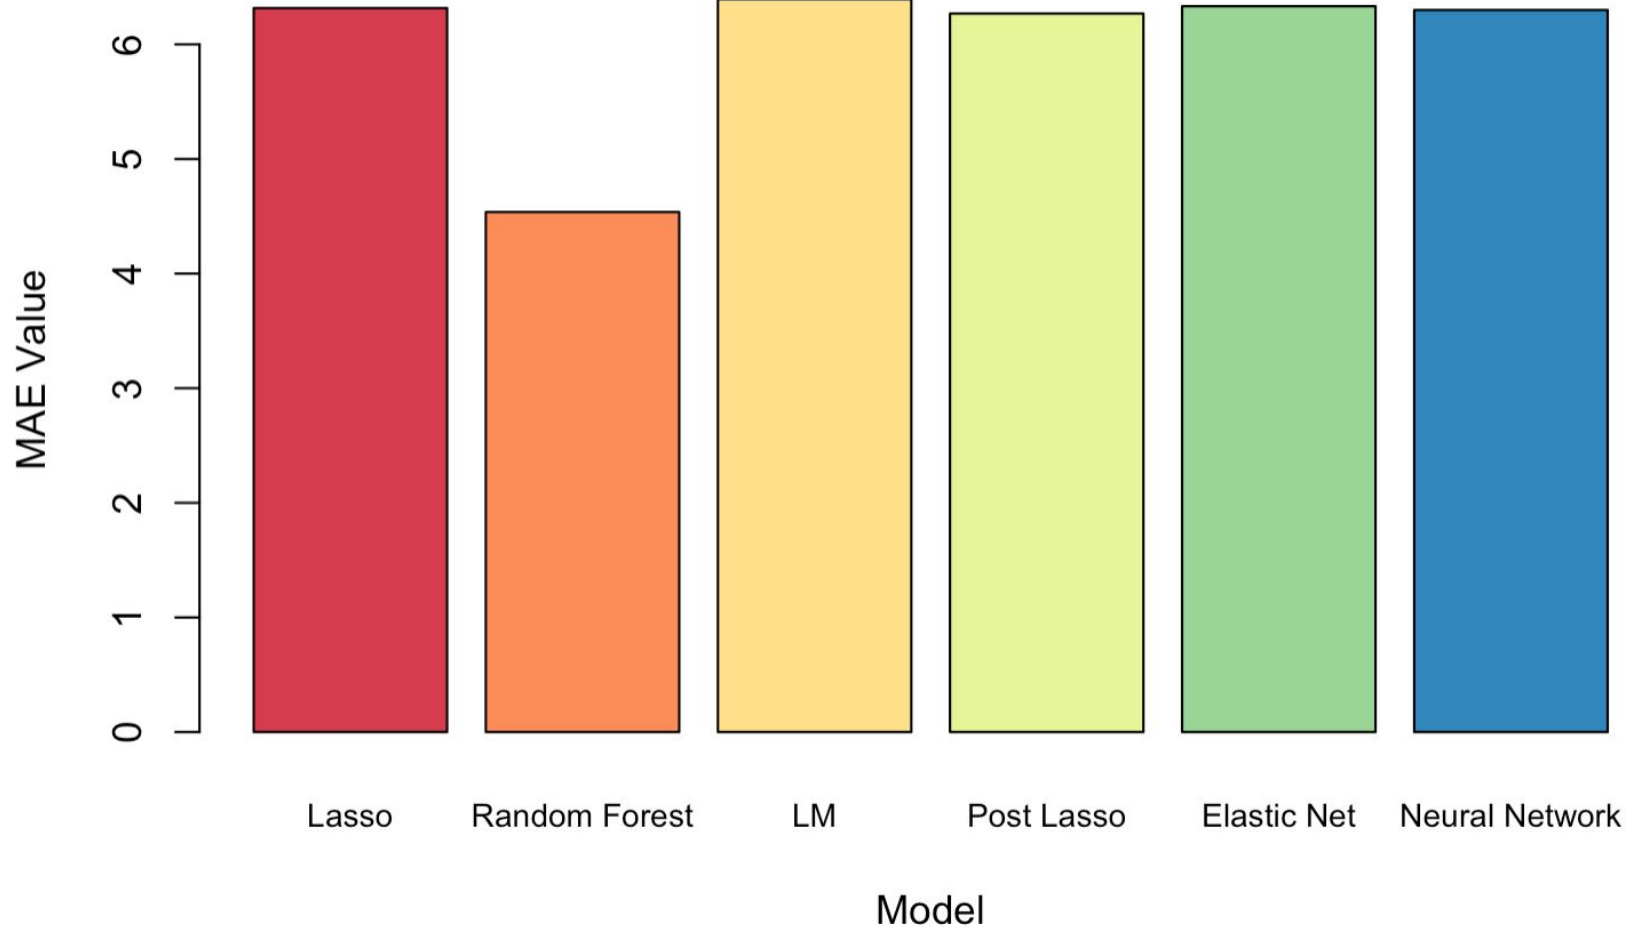
R2 by Model



RMSE by Model



MAE by Model





Conclusions

Conclusion

- ▷ We learned that a few variables matter a lot, and everything else doesn't really have an impact
- ▷ Management can maximize their market cap use by hiring someone highly skilled in a low-earning position that can help the team a lot but won't cost as much
- ▷ Teams with lower international reputations can hire highly skilled players just under a rating of 80

Thank you! Questions?

